

To print: Select File and then Print from your browser 's menu.

Embedded C in cars: Top bug-killing rules

Posted:20 Jan 2016

The [software](#) content of automobiles has grown phenomenally in the past few years. Estimates are that a modern luxury car may contain as many as 100 million lines of code. This code runs a host of interacting sub-systems from the infotainment systems, window controls, and air conditioning to the most critical engine control and safety systems. While many of these sub-systems are developed by the auto manufacturer, a good number of them come from third-party contractors as well as from independent aftermarket suppliers.

A number of standards are emerging in the auto industry such as the well established CAN bus, the AUTomotive Open System Architecture (AUTOSAR), the Media Oriented Systems Transport (MOST), and the ISO 26262 standard. While the first three standards address the external interaction and networking issues within automobiles, ISO 26262 is a set of requirements for safety and reliability. It is important for applications from different suppliers and groups within a manufacturer to conform to these standards, though only ISO 26262 calls for the use of well-defined coding rules to prevent bugs that can lurk within code and cause sometimes life-threatening malfunctions (**figures 1 and 2**).

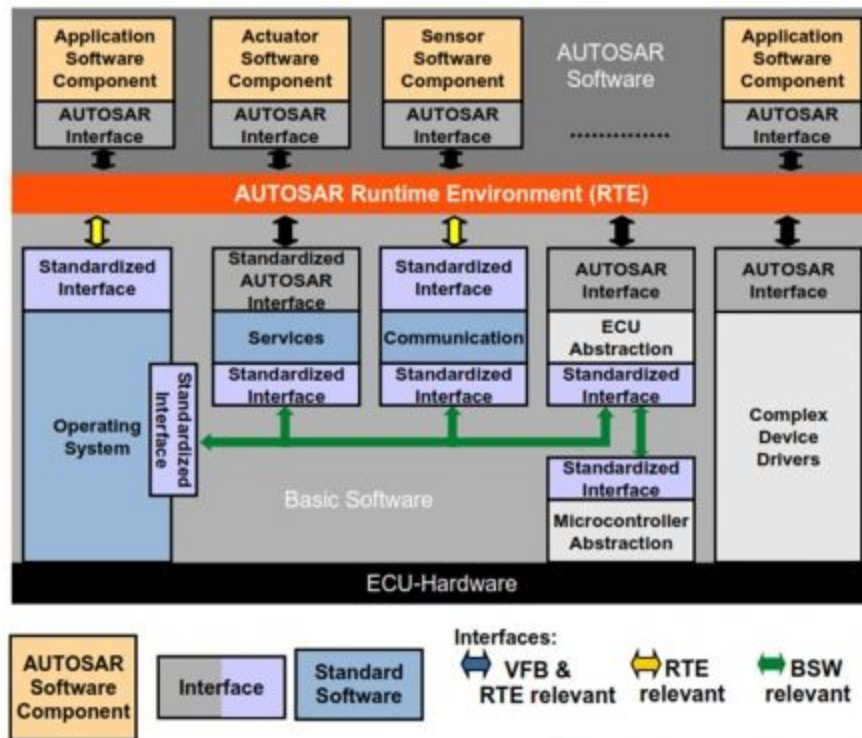


Figure 1: The AUTOSAR system links together various software applications running on different electronic control units (ECUs) within the automobile.

An embedded C coding standard is therefore needed that can be followed and understood by all teams and members of teams developing these interactive, networked systems. This is especially important because different teams and different suppliers will be using different development tools, compilers, and analysis tools. Therefore it is imperative to establish a common ground at the coding level.

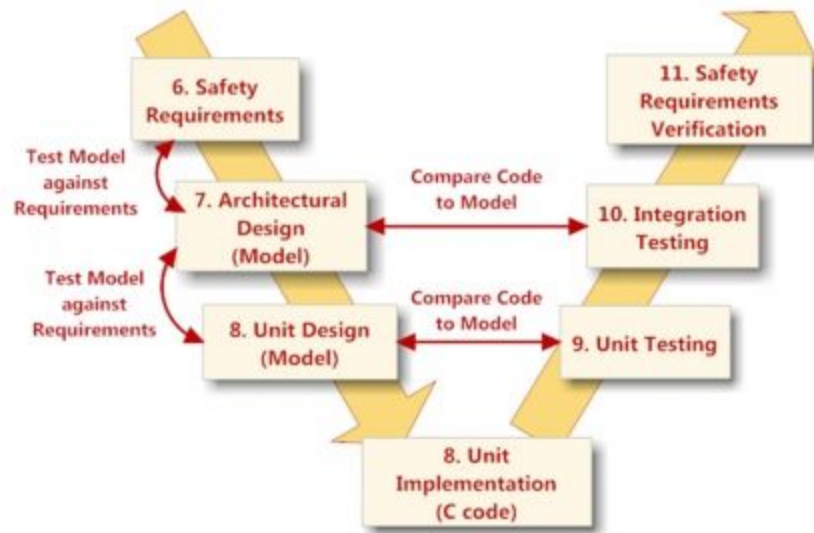


Figure 2: C coding is at the heart of ISO 26262 compliance but the standard does not specify coding rules or standards at the C level.

What's to gain from a coding standard?

The adoption of a coding standard by a team or a company has many benefits. For example, a coding standard increases the readability and portability of software, so that software may be maintained and reused at lower cost. A coding standard also benefits a team of software developers and their entire organisation by reducing the time required by individual team members to understand or review the work of peers.

However, one of the biggest potential benefits of a coding standard has been too long overlooked: a coding standard can help keep bugs out. It's cheaper and easier to prevent a bug from creeping into code than it is to find and kill it after it has entered. Thus, a key strategy for keeping the cost of firmware development down is to write code in which the compiler, linker, or a static-analysis tool can keep bugs out automatically—in other words, before the code is allowed to execute. While it is certainly important to use tools to verify and certify conformance to the standards mentioned above, such certification does not guarantee that the underlying code is bug-free.

That is because there are many sources of bugs in software programs. The original programmer creates some of the bugs, a few lurking in the dark shadows only to emerge months or years later. Additional bugs result from misunderstandings by those who later maintain, extend, port, and/or reuse the code.

The number and severity of bugs introduced by the original programmer can be reduced through disciplined conformance with certain coding practices, such as the placement of constants on the left side of each equivalence (==) test.

The original programmer can also influence the number and severity of bugs introduced by maintenance programmers. For example, appropriate use of portable fixed-width integer types (such as `int32_t`) ensures that no future port of the code to a new compiler or target processor will encounter an unexpected overflow.

The number and severity of bugs introduced by maintenance programmers can also be reduced through the disciplined use of consistent commenting and stylistic practices, so that everyone in an organisation can more easily understand the meaning and proper use of variables, functions, and modules.

Available C coding standards

Over the years, there have been many coding standards available, including several specifically aimed at the use of C for firmware development, as well as MISRA's "Guidelines for the Use of the C Language in Safety-Critical Systems." The authors of this so-called "MISRA-C" standard are knowledgeable of the risks in safety-critical system design and their guidelines narrow the C language to a safer subset. (For

more information about the MISRA-C Guidelines, see Nigel Jones's article "MISRA-C Guidelines for Safety Critical Software".)

Unfortunately, although firmware coding standards and MISRA-C sometimes overlap, coding standards too often focus primarily on stylistic preferences—missing their chance to help reduce bugs. MISRA-C, by contrast, offers great advice for eliminating bugs but very little guidance on practical day-to-day issues of style. And differences in style can lead to differences in understanding, which can introduce bugs.

To address this gap, the Barr Group Embedded C Coding Standard was developed. This coding standard was created from the ground up to help keep bugs out of firmware: each rule being optimised towards that objective. In addition, we applied the following guiding principles, which served to eliminate conflict over items that are sometimes viewed by individual team members as personal stylistic preferences:

- Individual programmers do not own the software they write. All software development is work for hire for an employer or a client and, thus, the end product should be constructed in a workmanlike manner.
- For better or worse (well, mostly worse), the ANSI/ISO "standard" C programming language allows for a significant amount of variability in the decisions made by compiler implementers. These many so-called "implementation-defined," "unspecified," and "undefined" behaviours, along with "locale-specific options", (see Appendix G of the ISO C standard), mean that programs compiled from identical C source code may behave very differently at run time. These grey areas in the language greatly reduce the portability of C programs that are not carefully crafted.
- The reliability and portability of code are more important than either execution efficiency or programmer convenience.
- The MISRA-C guidelines were carefully crafted and are worthy of study. On the small number of issues where we have a difference of opinion with MISRA-C, we make this clear. Of course, followers of Barr Group's coding standard may wish to adopt the other rules of MISRA-C in addition.
- To be effective in keeping out bugs, coding standards must be enforceable. Wherever two or more competing rules would be similarly able to prevent bugs but only one of those rules can be enforced automatically, the more enforceable rule is recommended.

Ten bug-killing rules for C

Here are some examples of coding rules you can follow to reduce or eliminate certain types of firmware bugs.

Rule #1—Braces

Braces { } shall always surround the blocks of code (also known as compound statements) following if, else, switch, while, do, and for keywords. Single statements and empty statements following these keywords shall also always be surrounded by braces.

```
// Don't do this ...
if (timer.done)
// A single statement needs braces!
timer.control = TIMER_RESTART;
// Do this ...
while (!timer.done)
{
// Even an empty statement should be surrounded by braces.
}
```

Reasoning: Considerable risk is associated with the presence of empty statements and single statements that are not surrounded by braces. Code constructs of this type are often associated with bugs when nearby code is changed or commented out. This type of bug is entirely eliminated by the consistent use of braces.

Rule #2—Keyword "const"

The const keyword shall be used whenever possible, including:

- To declare variables that should not be changed after initialisation,
- To define call-by-reference function parameters that should not be modified (for example, char const * p_data),
- To define fields in structs and unions that cannot be modified (such as in a struct overlay for memory-mapped I/O peripheral registers), and

- As a strongly typed alternative to `#define` for numerical constants.

Reasoning: The upside of using `const` as much as possible is compiler-enforced protection from unintended writes to data that should be read-only.

Rule #3—Keyword "static"

The `static` keyword shall be used to declare all functions and variables that do not need to be visible outside of the module in which they are declared.

Reasoning: C's `static` keyword has several meanings. At the module-level, global variables and functions declared `static` are protected from inadvertent access from other modules. Heavy-handed use of `static` in this way thus decreases coupling and furthers encapsulation.

Rule #4—Keyword "volatile"

The `volatile` keyword shall be used whenever appropriate, including:

- To declare a global variable accessible (by current use or scope) by any interrupt service routine,
- To declare a global variable accessible (by current use or scope) by two or more tasks, and
- To declare a pointer to a memory-mapped I/O peripheral register set (for example, `timer_t volatile * const p_timer`).

Reasoning: Proper use of `volatile` eliminates a whole class of difficult-to-detect bugs by preventing the compiler from making optimisations that would eliminate requested reads or writes to variables or registers that may be changed at any time by a parallel-running entity.

Rule #5—Comments

Comments shall neither be nested nor used to disable a block of code, even temporarily. To temporarily disable a block of code, use the preprocessor's conditional compilation feature (for example, `#if 0 ... #endif`).

```
// Don't do this ...
```

```
/*
```

```
a = a + 1;
```

```
/* comment */
```

```
b = b + 1;
```

```
*/
```

```
// Do this ...
```

```
#if 0
```

```
a = a + 1;
```

```
/* comment */
```

```
b = b + 1;
```

```
#endif
```

Reasoning: Nested comments and commented-out code both run the risk of allowing unexpected snippets of code to be compiled into the final executable.

Rule #6—Fixed-width Data Types

Whenever the width, in bits or bytes, of an integer value matters in the program, a fixed-width data type shall be used in place of `char`, `short`, `int`, `long`, or `long long`. The signed and unsigned fixed-width integer types shall be as shown in the **table**.

Standard for signed and unsigned fixed-width integer types.

Integer width	Signed type	Unsigned type
8 bits / 1 byte	<code>int8_t</code>	<code>uint8_t</code>
16 bits / 2 bytes	<code>int16_t</code>	<code>uint16_t</code>
32 bits / 4 bytes	<code>int32_t</code>	<code>uint32_t</code>
64 bits / 8 bytes	<code>int64_t</code>	<code>uint64_t</code>

Reasoning: The ISO C standard allows implementation-defined widths for char, short, int, long, and long long types, which leads to portability problems. Though the 1999 standard did not change this underlying issue, it did introduce the uniform type names shown in the table, which are defined in the new header file . These are the names to use even if you have to create the typedefs by hand.

Rule #7—Bit-wise Operators

None of the bit-wise operators (in other words, &, |, ~, ^, <<, and >>) shall be used to manipulate signed integer data.

// Don't do this ...

```
int8_t signed_data = -4;
```

```
signed_data >>= 1; // not necessarily -2
```

Reasoning: The C standard does not specify the underlying format of signed data (for example, 2's complement) and leaves the effect of some bit-wise operators to be defined by the compiler author.

Rule #8—Signed and Unsigned Integers

Signed integers shall not be combined with unsigned integers in comparisons or expressions. In support of this, decimal constants meant to be unsigned should be declared with a 'u' at the end.

// Don't do this ...

```
uint8_t a = 6u;
```

```
int8_t b = -9;
```

```
if (a + b < 4)
```

```
{
```

```
// This correct path should be executed
```

```
// if -9 + 6 were -3 < 4, as anticipated.
```

```
}
```

```
else
```

```
{
```

```
// This incorrect path is actually
```

```
// executed because -9 + 6 becomes
```

```
// (0x100-9) + 6 = 253.
```

```
}
```

Reasoning: Several details of the manipulation of binary data within signed integer containers are implementation-defined behaviours of the C standard. Additionally, the results of mixing signed and unsigned data can lead to data-dependent bugs.

Rule #9—Parameterized Macros vs. Inline Functions

Parameterized macros shall not be used if an inline function can be written to accomplish the same task.

// Don't do this ...

```
#define MAX(A, B) ((A) > (B) ? (A) : (B))
```

// ... if you can do this instead.

```
inline int max(int a, int b)
```

Reasoning: There are a lot of risks associated with the use of preprocessor #defines, and many of them relate to the creation of parameterized macros. The extensive use of parentheses (as shown in the example) is important, but doesn't eliminate the unintended double increment possibility of a call such as MAX(i++, j++). Other risks of macro misuse include comparison of signed and unsigned data or any test of floating-point data.

Rule #10—Comma Operator

The comma (,) operator shall not be used within variable declarations.

// Don't do this ...

```
char * x, y; // did you want y to be a pointer or not?
```

Reasoning: The cost of placing each declaration on a line of its own is low. By contrast, the risk that either the compiler or a maintainer will misunderstand your intentions is high.

Keeping bugs out of code requires constant vigilance and attention to detail. Following a set of consistent rules can help develop a mentality that will add to this success. Adding these ten easy-to-follow rules to your coding standard, and enforcing them, is a simple and effective way to keep large numbers of bugs

out of your C code.

About the author

Barr Group co-founder and Chief Technical Officer Michael Barr is a former adjunct professor of electrical and computer engineering with over a decade of software design and implementation experience. Internationally recognised as an expert in the field of embedded software process and architecture, Barr has been admitted as a testifying expert witness in U.S. and Canadian court cases involving issues of reverse engineering (including DMCA), interception of encrypted signals (Federal Communications Act), patent infringement, theft of copyrighted source code (including trade secrets issues), and product liability.

This article was printed from EE Times-Asia located at::

http://www.eetindia.co.in/ART_8800718970_1800001_TA_cdcd835f.HTM

[Back to Article](#) | [Back to Home](#)