

C++11 and Intel TBB

Klaas van Gend

October 26, 2012

Let me introduce myself

Name: Klaas van Gend

Affiliation: Vector Fabrics

Job title: Sr. FAE

History / Claim to fame:

- Started programming C at 14
- First experience with Linux in 1993
- Professional Linux dev since 1999
- Author of "umtsmon"
- Author of game "The Butterfly Effect"
- Author of approx 60 magazine articles

Hobbies:

- Glider piloting
- Card and Board games



Vector Fabrics

Founded: 2007

Motto: Do more with multicore!

Employees: 20

Address:

- Vector Fabrics B.V.
Vonderweg 22
5616 RM Eindhoven
The Netherlands
- +31 40 8200 960
- <http://www.vectorfabrics.com/>
- also on Twitter, LinkedIn and Facebook



C++ History

- 1979 Bjarne Stroustrup starts working on **C with Classes**
 - no true language, uses "C++ to C compiler"
- 1983 renamed to **C++**, version 1.0
 - adds virtual functions, operator overloading
- 1989 C++ 2.0 released
 - adds multiple inheritance, abstract classes,
 - adds static, const, protected members
- 1998 ISO Standard **C++98**
 - first official standard
 - adds templates, exceptions, namespaces, bool, STL
- 2003 ISO Standard C++03
 - only typos and inconsistencies fixed
- 2007 ISO Standard TR1
 - adds regular expressions, smart pointers, hash tables, random number generators
- 2011 ISO Standard C++0x turns **C++11**
 - adds *multithreading*, move semantics, lambdas

C++11 and gcc

With the latest compilers+libs, support is looking great!

Don't forget to specify `-std=gnu++11` or `-std=c++11`

- g++ 4.8 <http://gcc.gnu.org/projects/cxx0x.html>
 - nearing completion, missing are mostly low-level concurrency items
 - g++ 4.7 lacks TLS and alignment support
 - don't assume any compatibility with 4.5 and below
- Clang / LLVM 3.1 http://clang.llvm.org/cxx_status.html
 - nearing completion, missing are mostly low-level concurrency items
 - Clang 3.1 support is roughly comparable to GCC 4.7
- libstdc++
<http://gcc.gnu.org/onlinedocs/libstdc++/manual/status.html>
 - `thread::id` comparisons ill-defined - do not use
 - timed waits have minor issues
 - minor issues at thread exit

It's all in the libraries!

- Most of complexity and platform compatibility nicely hidden in SL and STL:
 - `std::thread` run lambda or function (object) as a thread
 - `std::mutex`, `std::lock_guard`, `std::lock` for locking
 - `std::unique_lock`, `std::defer_lock`, `std::adopt_lock` to manage locks
 - `std::condition_variable`
 - `std::atomic` templates for atomic access
 - barriers and fences
 - `std::futures`, `std::promises` nice higher level interface
 - `std::chrono` for accurate time-keeping
 - exception handling in threaded environments
- Do not forget to link against pthreads: `-lpthread`

std::thread example

- std::thread simplest use: function pointer and arguments
- return codes are **inaccessible!**

Listing 1: Hello Concurrent World

```
1 #include <iostream>
2 #include <thread>
3 using namespace std;
4
5 void hello(int n) {
6     cout << "Hello World " << n << endl;
7 }
8
9 int main(void) {
10     thread t(hello, 684);
11     t.join();
12     return 0;
13 }
```

std::thread more constructors

Listing 2: thread using lambda expression or function object

```
1 #include <iostream>
2 #include <thread>
3
4 class funobject
5 {
6 public:
7     void operator()(int n) const
8     {
9         std::cout << "Hi funobject " << n << std::endl;
10    }
11 };
12
13 int main(void)
14 {
15     funobject fo;
16     std::thread f(fo, 42);
17
18     std::thread l( [] (int n=684)
19         { std::cout << "Hi lambda " << n << std::endl; }
20     );
21
22     f.join();
23     l.join();
24 }
```

scope and lifespan

- `thread.join()` waits for a thread to complete
- `thread.detach()` detaches a thread from current scope
 - object no longer "joinable"!
 - thread object can go out of scope - doesn't matter
 - real (inaccessible) thread object destroyed when thread ends
 - access to variables doesn't change
- If a thread object goes out of scope and it wasn't joined or detached, **the application ends**.
- Transfer ownership using *move semantics* (new in C++11):

```
std::thread t1(some_function);  
std::thread t2 = std::move(t1);  
t1.join(); // error !
```

The four mutexes

`std::mutex`

- `std::mutex.lock()` locks the mutex, blocks if already locked
- `std::mutex.try_lock()` returns true on success or false if already locked
- `std::mutex.unlock()` unlocks

`std::timed_mutex` adds:

- `std::timed_mutex.try_lock_for(d)` waits max duration d
- `std::timed_mutex.try_lock_until(t)` waits until time t

Also available:

- `std::recursive_mutex` allows multiple locks in same thread
- `std::recursive_timed_mutex`

A better locking mechanism

- Hard to do exception safe locking using `.lock()` / `.unlock()`
- Better: use `std::lock_guard`
 - its constructor locks, its destructor unlocks
 - as long as the object is in scope, your lock is held

Listing 3: using `std::lock_guard`

```
1 int theI = 0;
2 std::mutex theMutex;
3
4 void safe_inc(int n)
5 {
6     std::lock_guard<std::mutex> lock(theMutex);
7     theI += n;
8     std::cout << "now: " << theI << std::endl;
9 }
10
11 int main()
12 {
13     std::thread t1(safe_inc, 42);
14     std::thread t2(safe_inc, 684);
15
16     t1.join();
17     t2.join();
18 }
```

std::unique_lock

- std::lock_guard locks the mutex upon creation
- If you need more flexibility or need to it move around:
unique_lock<mutex> myObject(aMutex, adopt_lock)
 - Operates just like std::lock_guard
 - But consumes slightly more memory and cycles
- unique_lock<mutex> myObject(aMutex, defer_lock)
 - Creates myObject, but **doesn't** lock aMutex yet
 - Convenient if you need to lock two locks at the same time using e.g. std::lock()
- You use std::unique_lock like a local mutex
 - But still have exception safeness...
 - And you can std::move() the std::unique_lock !!!

Atoms in C++

`std::atomic<type>`

- Where type can be any type.
- Is a (templated) class - and thus is a type in itself
- Class members include:
 - `.load()`
 - `.store()`
 - `.exchange()`
 - `.compare_exchange_weak()` and `.compare_exchange_strong()`
- If type is an integral type like `int` or `double`, there are additional members:
 - `.fetch_add()`, `.fetch_sub()`, `.fetch_or()`, `.fetch_and()`, `.fetch_xor()`
 - each of the above also exists as `+=` and similar
 - there are also overloaded `++` and `-` operators

C++ atomics in practice

The easy part:

<code>int myValue = 684;</code>	not atomic
<code>myValue++;</code>	not atomic
<code>atomic<int> myAtom = myValue;</code>	assignment works
<code>myAtom += 42;</code>	atomic addition
<code>myValue = myAtom.fetch_sub(3);</code>	atomic subtraction
Both variables will now contain $684 + 42 - 3$	
Note that myValue did not become atomic!	

C++ also allows different memory synchronization mechanisms.

- a topic that gives you a headache, guaranteed!
- suggested reading: *C++ Concurrency in Action* by Anthony Williams, chapter 5.

Futures and Promises

- The work flow so far:
 - use `std::thread` to spawn a new thread
 - make sure all data is available
 - take care of protecting read/write access to shared data
 - `.join()` when you need the output of the thread
 - and make sure the output is still available at that stage
- Every step complicated due to possible exceptions
- By now, everyone is planning their own thread class abstraction?
- Hopefully not!

Example using std::async and std::future

Listing 4: workload: reverse string

```
1 #include <string.h>
2
3 char* revString(char* arg)
4 {
5     // WARNING mem leak here!!!
6     char* result = strdup(arg);
7     size_t end = strlen(arg);
8     for (int i=0; i< end; i++)
9         result[i] = arg[end-i-1];
10    return result;
11 }
```

Listing 5: serial version

```
13 int smain(int argc, char* argv[])
14 {
15     for(int i=1; i<argc; i++)
16         cout <<revString(argv[i]) <<" ";
17     cout << endl;
18 }
```

Listing 6: parallel version

```
1 #include <future>
2
3 int fmain(int argc, char* argv[])
4 {
5     std::future<char*> myV[argc];
6
7     for(int i=1; i<argc; i++)
8         myV[i-1] =
9             std::async(revString, argv[i]);
10
11     for(int i=0; i<argc-1; i++)
12         cout << myV[i].get() << " ";
13     cout << endl;
14 }
```


std::async and std::future - 1

- `std::async()` has similar notation to `std::thread()` :
 - `std::async(myFunction)`
 - `std::async(myFunction, variable)`
 - `std::async(&Class::member, &instance, var)`
- Conceptually, they are different:
 - a `async` starts an asynchronous function call
 - a `thread` is something that runs in parallel to the current thread.
 - so an `async` *could* be implemented as a separate thread

std::async and std::future - 2

- a future holds a value that should become available sometime as the result of an asynchronous function call.
- `std::future<T>` holds the *future* return value of type T
- Calling `std::async`, running the asynchronous code can happen:
 - immediately after start
 - only once you call the `.get()`
 - you can specify by adding another parameter to `std::async` :
 - `async(std::launch::async|std::launch::deferred, myFunction)`
 - (the default behavior specifies both)
- Exceptions thrown inside `myFunction()` will be caught, kept and rethrown in the `.get()`

std::packaged_task

- std::async has a few limitations
 - no control over the wrapped "task"
 - future is returned immediately - and is only movable.
- Higher level abstraction: the std::packaged_task template.
- Example notation:
 - std::packaged_task<int()> task([]() {return 7;});
 - <int()> implies a function object that returns an int
 - []() {return 7;} is a lambda that just returns 7.
- The task is now wrapped, can be executed anywhere anytime.
- A future return value can be obtained via .get_future()
- The wrapped task (lambda in our case) will be executed:
 - *synchronously* if we call the operator() on the packaged_task
 - *asynchronous* if we std::move the packaged_task into a thread
 - *"likely" asynchronous* if we just call .get() on the future

std::promise

- A `std::promise` is a vehicle to move parameters between threads
- It holds all necessary synchronization primitives
- Used internally by `std::packaged_task`
 - It passes the arguments residing in the `packaged_task` to the wrapped function
 - It hides the `std::future` from the wrapped function, but puts the return value in the future
 - It helps in catching all exceptions and storing them in the future
- "Fun":
 - `std::move` a `std::promise`
 - request its future through `.get_future()`
 - an exception of type `std::future_errc::broken_promise`.
- You probably don't need to toy with promises yourself.

Shared futures

- `std::future` is only `std::moveable`
- But what if you need the answer to a certain operation for two other operations?
 - Use a `std::shared_future` !
- This enables all kinds of 'fun' coding, like functional programming.
- No need for explicit locks to have e.g. spreadsheet cells depend on each other
- The *language* will sort out the ordering in which futures can be calculated!
- But let's not dive into that here...

Waiting...

- `future.get()` waits indefinitely
- But what if you don't want to wait that long?
 - In addition to `.lock()` `.get()` and C++11 defines things like `.wait_until()`, `.wait_for()` and `.try_lock_until()`
- Both require timing notation: a *point in time* and *duration*
 - neither being defined well in C
 - nor POSIX
 - let alone across platforms
- Cue the chrono library:
 - `std::time_point`, a point in time
 - `std::duration`, a time interval
 - `std::system_clock`, wall clock time from the system-wide realtime clock
 - `std::steady_clock`, monotonic clock that will never be adjusted
 - `std::high_resolution_clock`, the clock with the shortest tick period available

Summary

C++11

- Adds several layers of multithreading concepts:
 - "Basic":
 - `std::thread`
 - `std::mutex` and `std::lock_guard`
 - `std::atomic`
 - "Abstracted":
 - `std::packaged_task` and `std::async`
 - `std::future` and `std::promise`
- It also provides abstract but accurate timing and waiting
- Now only we have to wait for 100% compliant compilers and libraries...

Parallelize an encheferizer

Encheferizer

- Also known as *Jive filter*
- Converts text into "Svedeesh", Bork Bork Bork.
- Example code written in 'normal' C++
- Your jobs:
 - Parallelize the code using C++11 constructs
 - Prove it is paralellized
 - Prove it still has the same output



Svedeesh Cheff

Frum Veekipedia, zee free-a incyclupedeea

Zee **Svedeesh Cheff** is a [Mooppet](#) vhu eppeered in zee lung-roonneeng [Zee Mooppet Shoo](#) und ves oopereted by [Jeem Hensun](#) und [Frunk Ooz](#) seemooltuneuoosly. Bork Bork Bork!

Cuntents [\[hide\]](#)

[1 Cherecter](#)

Listing 7: serial encheferizer main

```
1 #include <cstdio>
2 #include <fstream>
3 #include <sstream>
4 #include <iostream>
5
6 extern std::string chefLine(const std::string& aLine);
7
8 int smain(void)
9 {
10     std::ifstream file("infile.txt");
11     char line[512];
12     while(!file.eof())
13     {
14         file.getline(line, 511);
15         std::cout << " " << chefLine(line) << std::endl;
16     }
17 }
```

Parallelize an encheferizer

Solution

Important things to remember:

- Do not forget to link against pthread: `-lpthread` !!!
- `std::async` needs to always have access to its arguments
- Split the `while` loop:
 - reading the file and starting the asyncs
 - `.get()`ing the futures
- Do not forget to force `async` to actually work asynchronously

Solution

Listing 8: parallel encheferizer main

```
1 #include <fstream>
2 #include <iostream>
3 #include <future>
4 #include <vector>
5
6 extern std::string chefLine(const std::string& aLine);
7
8 int pmain(void)
9 {
10     std::vector<std::string> myLines;
11     std::vector<std::future<std::string>> myConvertedLines;
12
13     std::ifstream file("infile.txt");
14     char line[512];
15     while(!file.eof())
16     {
17         file.getline(line, 511);
18         myLines.push_back(line);
19         myConvertedLines.push_back(
20             std::async(std::launch::async, chefLine, myLines.back()) );
21     }
22
23     // note the &: myLineFuture is not the iterator but the future
24     for(auto& myLineFuture : myConvertedLines)
25         std::cout << " " << myLineFuture.get() << std::endl;
26 }
```

Is it really multithreaded?

Listing 9: parallel encheferizer main

```
1 > gdb ./main.out
2 GNU gdb (GDB) SUSE (7.3-41.1.2)
3 ...
4 Reading symbols from main.out...done.
5 (gdb) set args p
6 (gdb) run
7 Starting program: main.out p
8 [New Thread 0x7ffff70b7700 (LWP 7255)]
9 [New Thread 0x7ffff68b6700 (LWP 7256)]
10 [New Thread 0x7ffff60b5700 (LWP 7257)]
11 [Thread 0x7ffff70b7700 (LWP 7255) exited]
12 [Thread 0x7ffff60b5700 (LWP 7257) exited]
13 [New Thread 0x7ffffefff700 (LWP 7258)]
14 [New Thread 0x7ffffef7fe700 (LWP 7259)]
15 [Thread 0x7ffff68b6700 (LWP 7256) exited]
16 [New Thread 0x7ffffeeffd700 (LWP 7260)]
17 [Thread 0x7ffffefff700 (LWP 7258) exited]
18 [New Thread 0x7ffffee7fc700 (LWP 7261)]
19 [Thread 0x7ffffef7fe700 (LWP 7259) exited]
20 [Thread 0x7ffffeeffd700 (LWP 7260) exited]
21 [New Thread 0x7ffffedffb700 (LWP 7262)]
22 [New Thread 0x7ffffed7fa700 (LWP 7263)]
23 [Thread 0x7ffffee7fc700 (LWP 7261) exited]
24 [Thread 0x7ffffedffb700 (LWP 7262) exited]
25 [Thread 0x7ffffed7fa700 (LWP 7263) exited]
26 Threedeeng fecileetees
27 ...
```

Questions?

Questions, anyone?

klaas AT vectorfabrics.com
<http://www.vectorfabrics.com/>